

SELF-DRIVING CARS

PERCEPTION




Photogrammetry & Robotics Lab

Perception for Self-Driving Cars Vision-based Approaches


Jens Behley

Part of the Course: Techniques for Self-Driving Cars by
C. Stachniss, J. Behley, N. Chebrolu, B. Mersch, L. Peters, I. Bogoslavskyi


Which information is needed?



Which information is needed?



Content of this lecture



- Overview of perception stack & tasks
- Dive into camera-based perception

Perception Suite



- Common Perception Suite
 - (Stereo) Camera
 - LiDAR
 - Radar

(Stereo) Camera



- Pro: cheap, high resolution, color
- Con: strongly affected by illumination, needs additional light at night

LiDAR Sensors



- Pro: Independent of illumination, Precise distance measurements
- Con: Expensive, mid resolution

Radar



- Pro: Position + velocity information, Matured technology
- Con: low resolution/sparse

Other Sensors

- Ultrasonic (near range)
- GPS
- Inertial Measurement Unit (IMU)
- Odometer

Complementing Modalities



- Many sensors are already build into cars (camera, RADAR, ultra-sound sensors, ...)
- Not a single sensor will enable self-driving, but combination of sensors
- **But: Sensor fusion** essential to integrate information from multiple sensors

Perception Suite



■ Common Perception Suite

- (Stereo) Camera

- LiDAR

- Radar

- GPS+IMU


Perception Tasks



Classification



Semantic Segmentation




Detection



Panoptic Segmentation

Convolution Neural Network



- Since 2012 success in ImageNet Challenge, the basis for most image-based perception tasks nowadays
- Here: high-level overview of CNNs

Why are CNNs successful now?

- Several reasons made progress possible:
 1. Availability of large-scale data (ImageNet, etc.)
 2. Availability of compute capabilities (GPUs)
 3. **Availability of code (and frameworks)!**
- Implementation for most paper available
- Many frameworks made it simple to build and train networks (Caffe, Theano, Torch, etc.)

Deep Learning Frameworks

- All operations must be implemented using GPU
- DL Frameworks available implementing the aforementioned operations (and many more)



TensorFlow



PyTorch

Convolution Neural Network



*

$W_{0,0}$	$W_{0,1}$	$W_{0,2}$
$W_{1,0}$	$W_{1,1}$	$W_{1,2}$
$W_{2,0}$	$W_{2,1}$	$W_{2,2}$

=

-77	-71	277
127	-95	87


$$\sum I(x + u, y + v)K(u, v)$$

- Convolution “slides” kernel/filter K over image I
- Trivia: Most DL frameworks use cross-correlation instead

Convolution Neural Network



$$H \times W$$



$$3 \times 3 \times 4$$



$$H \times W \times 1$$


(with zero padding)

- For multi-channel input convolutional kernel has also as many channels
- Produces still one activation map

Channel-first vs. Channel-last


- Organizing tensors in different ways possible
- Main conventions:
 - Channel-first (PyTorch): $C \times H \times W$
 - Channel-last (Tensorflow): $H \times W \times C$
- We stick now to the channel-first convention...

Convolutional Layer





- Use multiple kernels to produce C_{out} maps

ConvLayer + Activation Function




- Activation function (such as ReLU) applied after each convolutional layer
- Usually only implicit in the graphical representation

Activation Function




- Non-linear Activation functions
- Popular: **Rectified Linear Unit (ReLU)**
+ variants, e.g., Leaky ReLU

Convolution Neural Network




- Stack of convolutional layers
- Pooling layer to increase receptive field of layers

Receptive field



- Location in deeper layers take inputs of window of earlier layers
- Deeper layers “see” more from earlier layers

Pooling Layer



- Pooling layers increase the receptive field & aggregates information
- Translation invariance to small shifts
- Common: max pooling, average pooling

Example: Max Pooling

12	14	1	4	4	1
3	4	5	2	2	3
8	9	12	3	4	7
8	3	4	3	3	4

14	5	4
9	12	7

2 × 2 max pooling, stride 2

- Compute maximum in each region

Strided Convolution

$S = 1$

H



H



$S = 2$

H




H/2



W/2


- Common: Use **strided convolution** for downsampling to increase the receptive field
- Stride > 1 reduces size of feature map

Convolution Neural Network



- Output usually appropriately shaped tensor
- Example: C classes \rightarrow C logit values for softmax

Fully connected (FC) layer



- Each value of the input is used to produce output value: $y = \mathbf{W}\mathbf{x}$
- Common: flattening of $C \times H \times W$ tensor to vector $C \cdot H \cdot W$ before FC layer
- Also called **linear layer**

Learning

- Neural network is basically just a rather complex function:

$$f(\mathbf{x}_i; \theta) = L_3(L_2(L_1(\mathbf{x}_i; \mathbf{w}_1); \mathbf{w}_2); \mathbf{w}_3)$$

- **Loss function** is the objective we want to minimize; determines what network structure should learn

Common Loss Functions

- Loss $\ell(y_i, f(\mathbf{x}_i; \theta)) \in \mathbb{R}$ determines difference between prediction $f(\mathbf{x}_i; \theta)$ and target y_i

- Examples:


- **L2 loss** for regression tasks

$$\ell(y_i, f(\mathbf{x}_i; \theta)) = (y_i - f(\mathbf{x}_i; \theta))^2$$

- **Cross entropy** loss for classification

$$\begin{aligned}\ell(j, f(\mathbf{x})) &= -\log \frac{\exp(f_j(\mathbf{x}))}{\sum_k \exp(f_k(\mathbf{x}))} \quad \leftarrow \text{Softmax} \\ &= -f_j(\mathbf{x}) + \log(\sum_k \exp(f_k(\mathbf{x})))\end{aligned}$$

Learning via gradient descent




- **Idea:** Determine how to change parameters in a layer to reduce loss
- Parameter updates efficiently computed via **back propagation**

More details on CNNs

- We touched only parts needed to understand papers in the seminar
- Much more theory, building blocks, best practices
- For more details on CNNs, see the links in the description of the video.


Perception Tasks



Classification



Semantic Segmentation



Detection



Panoptic Segmentation

Object detection task




- Input: RGB Image
- Output:
 - **bounding boxes** defined by (x, y, w, h) or (c_x, c_y, w, h) or (x_1, y_1, x_2, y_2)
 - **confidence scores** in $[0, 1]$



Modern Object Detectors

- Rely mainly on Convolution Neural Networks (CNN)
- Two main paradigms:
 - **Anchor-based** approaches
 - **Anchor-free** approaches

Anchor-based Approaches



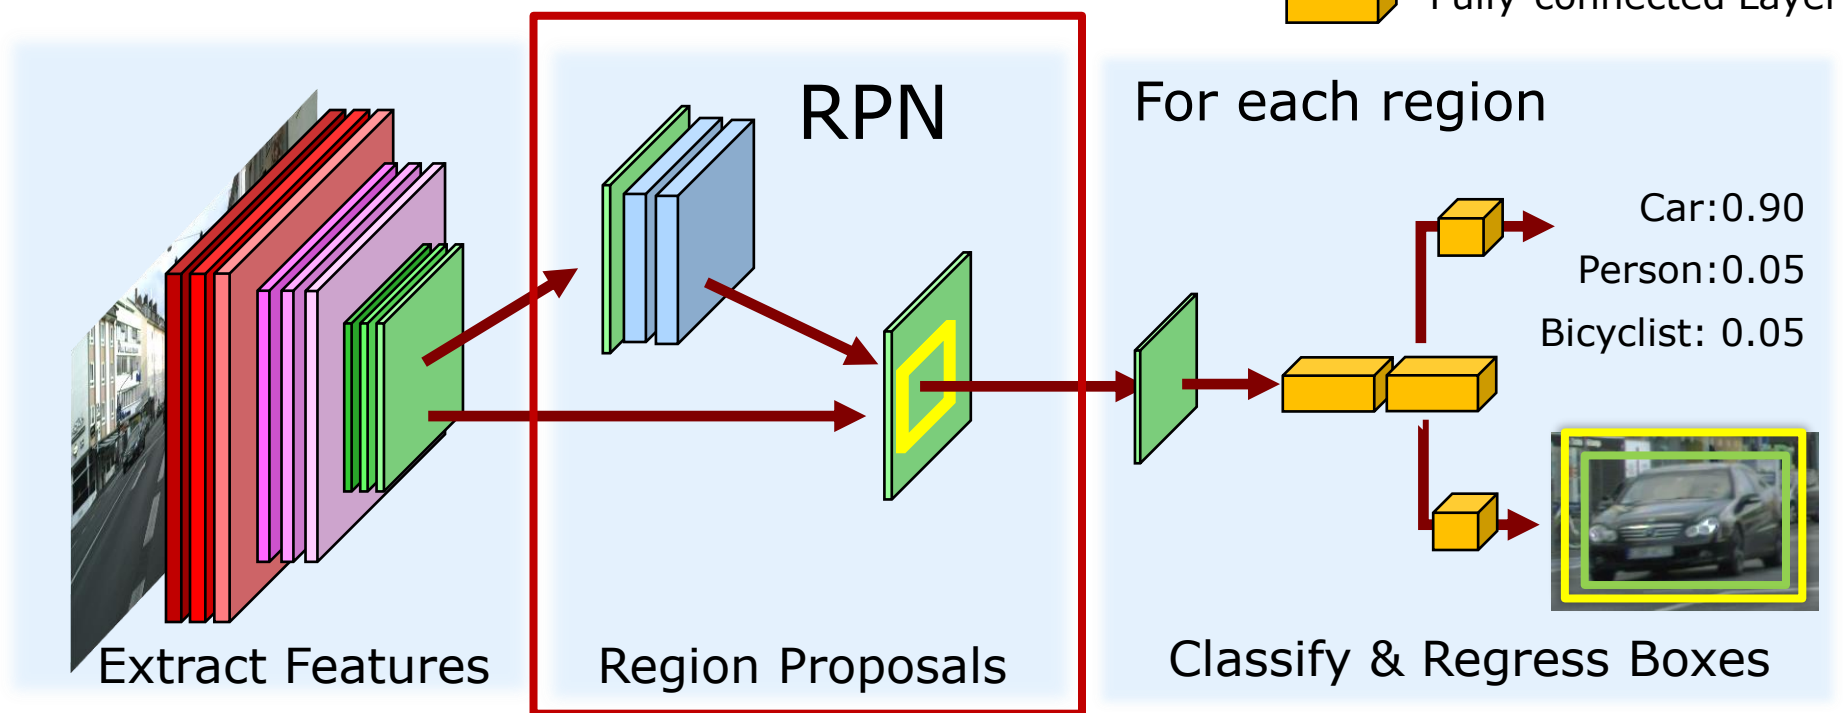
Two-stage



Single-stage


- **Two-stage** approaches:
 - First stage: Generate proposals based on anchors
 - Second stage: Refine and classify proposals
- **Single-stage** approaches:
 - Refine & classify anchors in single pass

Faster R-CNN (2015)



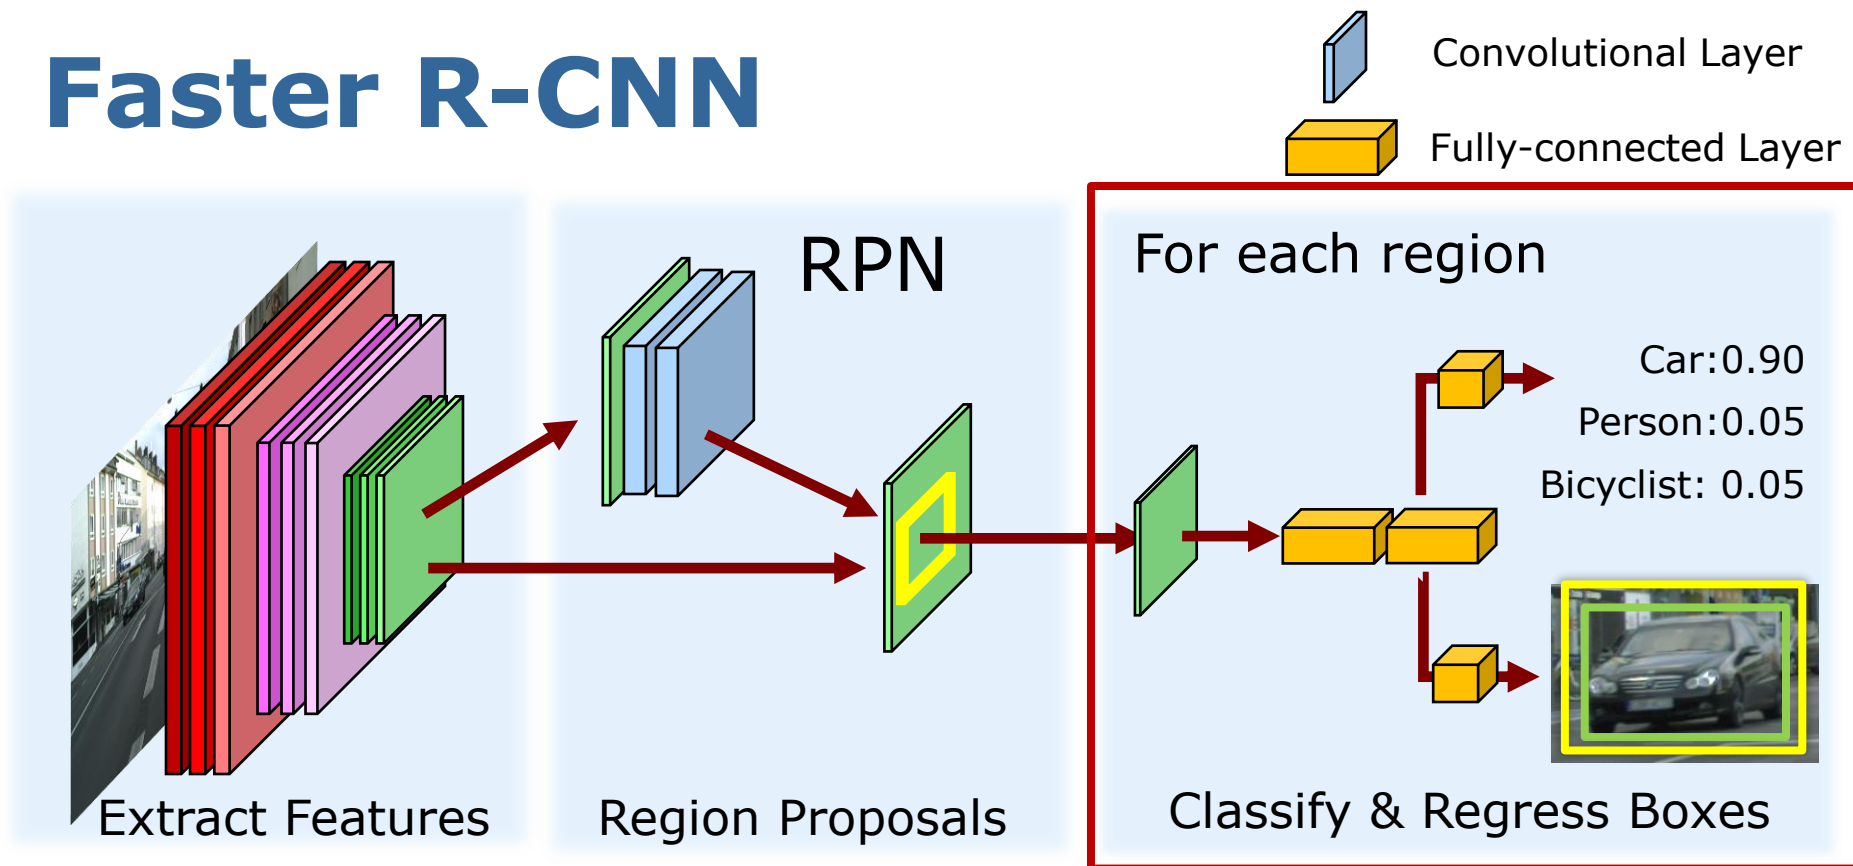
- **Region Proposal Network** produces proposals
- **Region-wise classification network** uses same features as input as RPN (shared features)

Region Proposal Network (RPN)




- Scores set of anchors with fixed initial sizes
- Produces objectness and anchor offsets (for each anchor)
- Keep **N-top scored anchors** as RoI for classification in second stage

Faster R-CNN



- **Region-wise classification network** classifies and refined bounding boxes (regression)
- But, how to extract proposal-specific features?

Region-of-Interest(RoI) Pooling




- Adaptive max pooling brings extracted feature maps into appropriate size for RoI Network

Faster R-CNN Summary

- Jointly trainable RPN and RoI classifier
→ sharing of features possible
(alternated training)
- Fast enough for near real-time
operation (~ 10 Hz)
- RPN already provides object bounding
boxes → second stage needed?

You Only Look Once (YOLO)

$$S = 4$$



Per anchor scores & offsets:

$$(O_1, \delta x_1, \delta y_1, \delta w_1, \delta h_1)$$

$$(O_2, \delta x_2, \delta y_2, \delta w_2, \delta h_2)$$

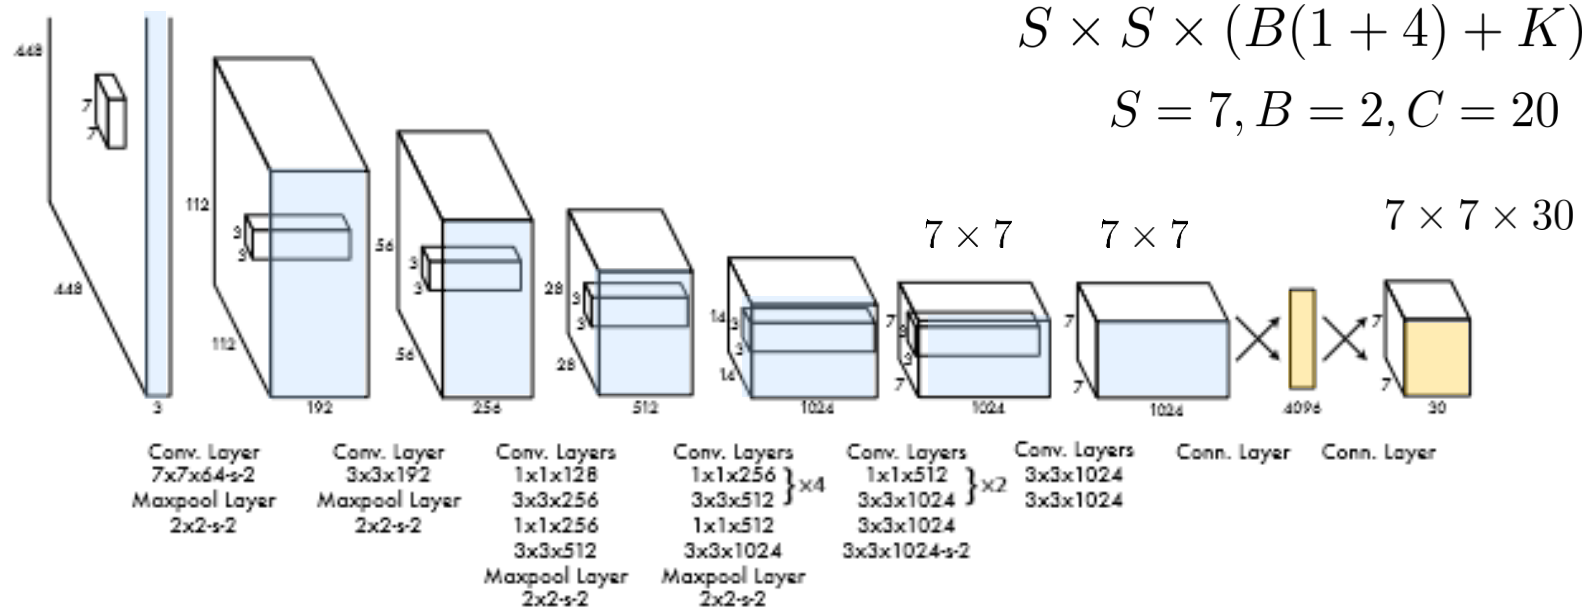
Per cell class scores:

$$(C_1, C_2, \dots, C_K)$$

$$\text{Output: } S \times S \times (B(1 + 4) + K)$$

- Predicts bounding boxes with single forward pass
- Each anchor gets objectness score O , bounding box offsets
- Objectness + class score determines outcome

YOLO Architecture




- YOLO uses a 24-layer convolutional network (DarkNet) with 2 fully connected layers, 2 anchors, 20 classes (Pascal VOC)

YOLO Summary


- Each grid cell produces at most 1 bounding box
- Only single pass through CNN needed
→ Blazingly fast (up to 144 Hz)
- But less proposals than Faster R-CNN
→ less accurate, misses too close objects

Anchor-free Approaches



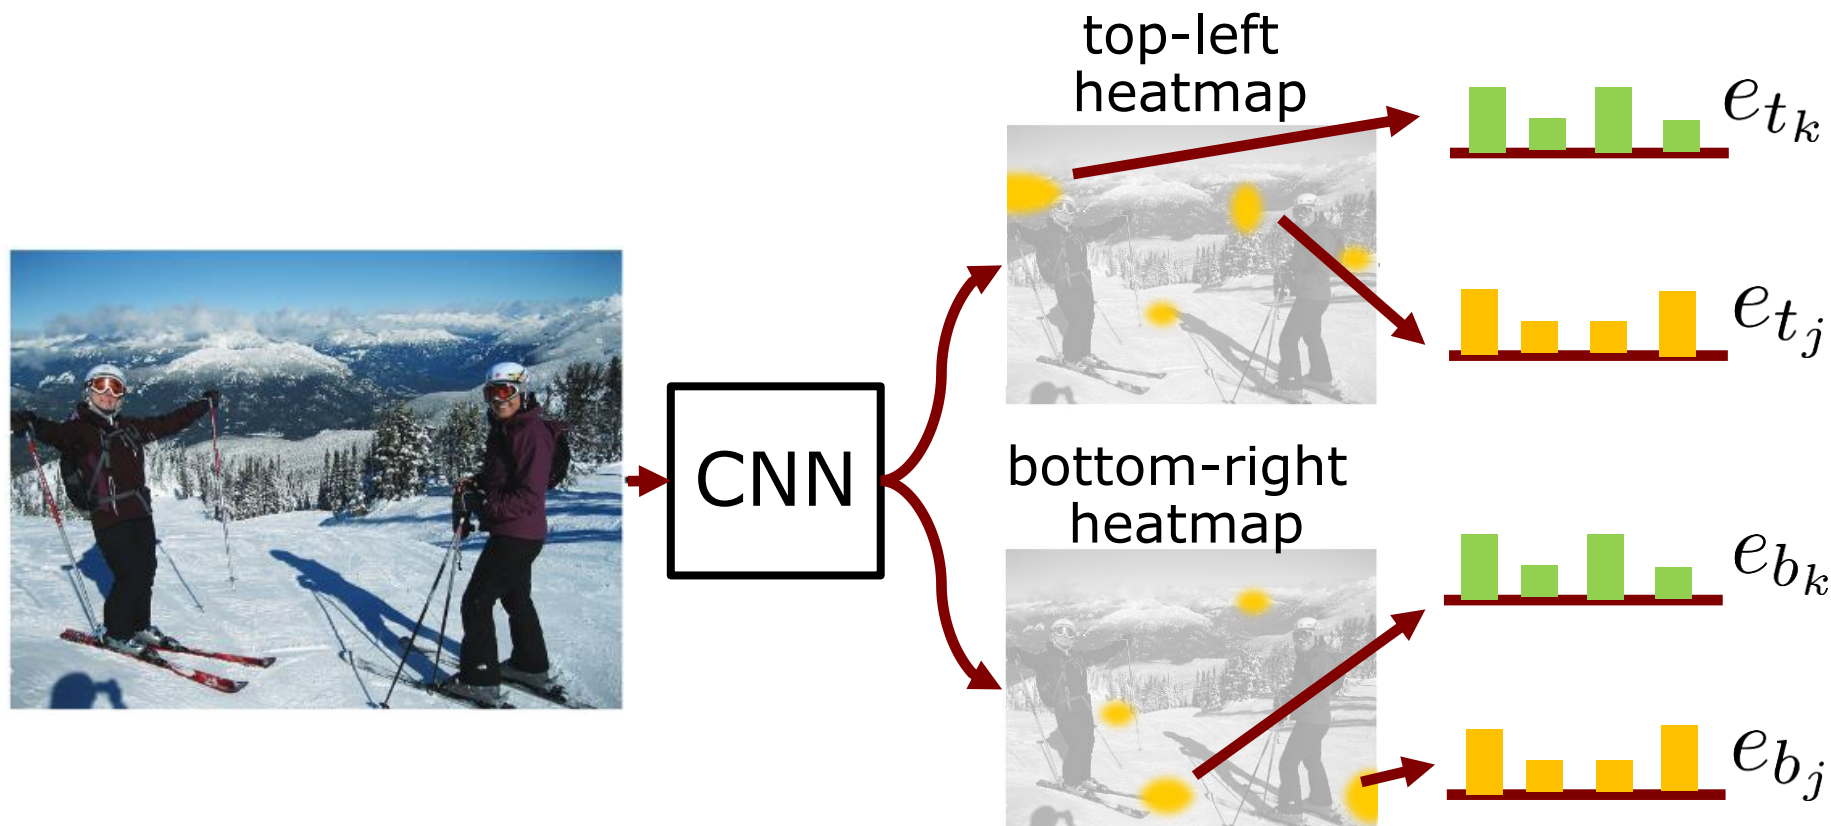
- Determine keypoints (e.g., corners, centers of bounding boxes)
- Post-processing uses keypoints to produce bounding boxes

Bounding boxes from corners



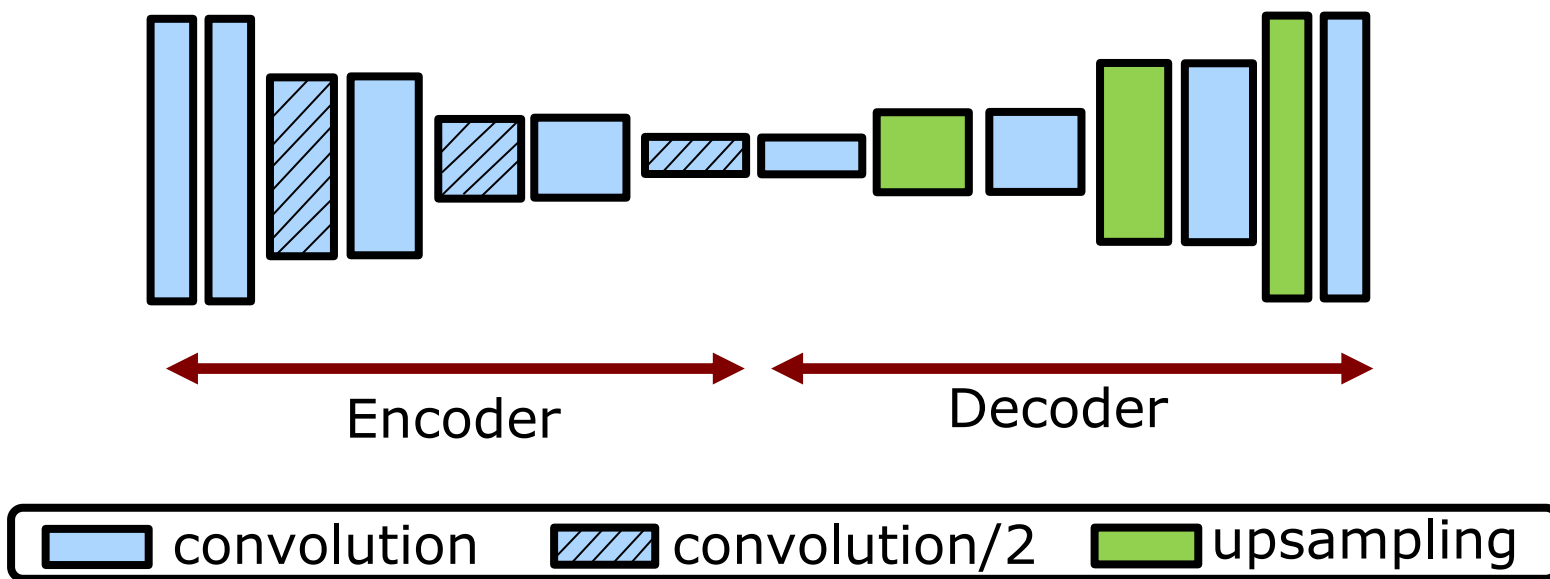
- Instead of scoring of anchors, **CornerNet** determines corners of bounding boxes
- Produce heatmaps of likelihood that at given pixel is upper-left or bottom-right corner

How to associate corners?




- To associate top-left and bottom-right corners, CornerNet determine **embedding** (“features”)
- Similar embeddings correspond to the same object

Encoder Decoder Architecture




- We want to have **pixel-wise** features
- Encoder uses **strided convolutions** or **max pooling** to down-sample feature maps
- Decoder upsamples feature maps to original resolution using **upsampling** operations

Common Upsampling Methods



- **Nearest neighbor upsampling** just copies values from nearby pixels

Common Upsampling Methods



- Learnable weights for interpolation:
Transpose Convolutions “inverts” convolution₅₁

CornerNet Summary

- **Anchor-free** approach
 - Corner locations (upper-left, bottom-right) with embedding vectors
 - Similarity between learned embedding vectors determines bounding box
- Encoder-Decoder architecture produces pixel-wise outputs


Perception Tasks



Classification



Semantic Segmentation



Detection



Panoptic Segmentation

Semantic Segmentation



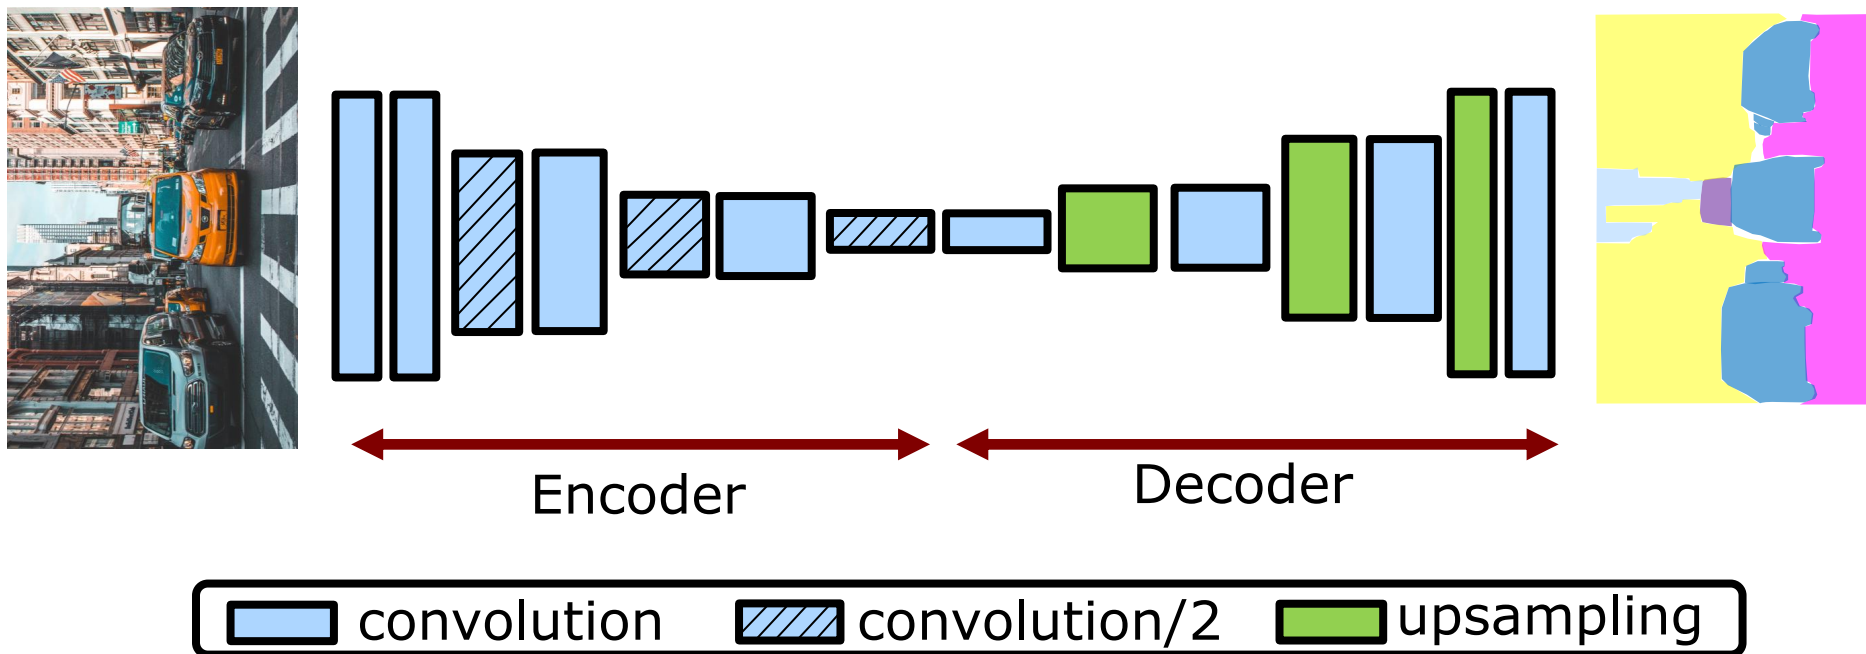
CNN



■ Building ■ Road ■ Car ■ Bus

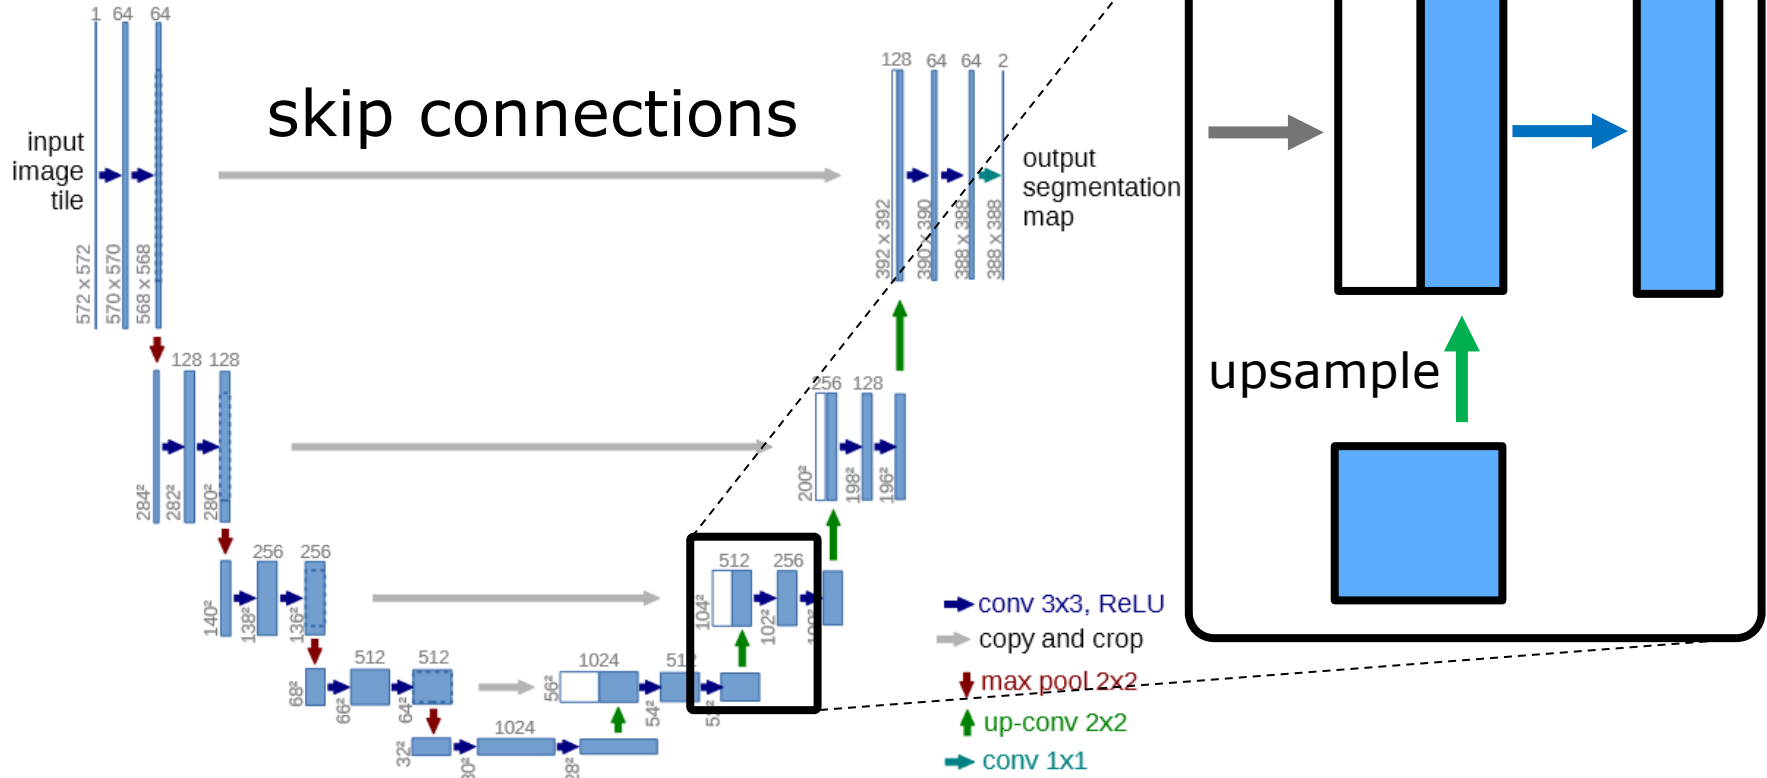
- **Goal:** Provide label $y_{i,j} \in \{1, \dots, K\}$ for each pixel in the image.

Encoder-Decoder Architecture



- Combine up-sampling with convolutional layers to regain spatial resolution

U-Net



- **Skip connections** help to retain fine-grained results
- Concatenate feature volumes from encoder and upsampled feature volumes from decoder
- Convolve to reduce number of channels

U-Net Summary

- **Encoder-decoder** architecture to produce pixel-wise logits for classification
- **Skip connections** to use high-resolution information from encoder


Perception Tasks



Classification



Semantic Segmentation




Detection




Panoptic Segmentation

Panoptic Segmentation




- Panoptic Segmentation unifies **semantic** and **instance** segmentation
- Distinguish **stuff** (e.g., vegetation, road, ...) and **thing** classes (e.g., car, pedestrian, ...)

Top-down vs. Bottom-up Instance Segmentation



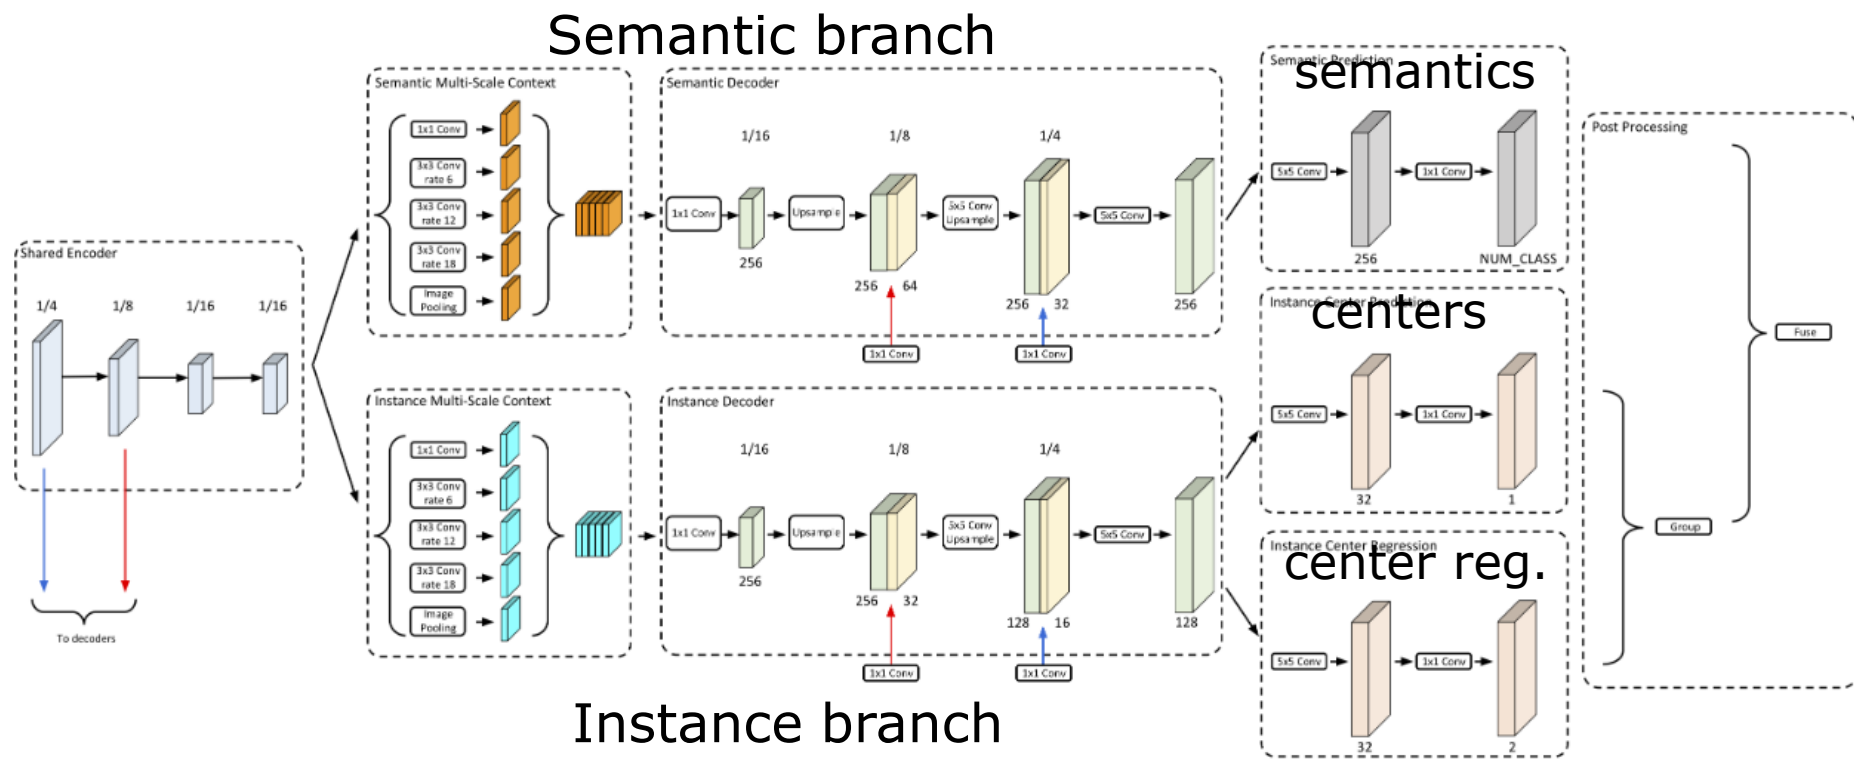
Top-down Approach



Bottom-up Approach

- **Top-down**: Instances are first determined and then foreground/background mask estimated
- **Bottom-up**: Determine per-pixel properties that are then used to cluster instances

Panoptic-DeepLab



- **Bottom-up approach** using separate branches for semantic and instance segmentation
- Use semantic labels to filter instances & majority vote on instances to assign instance labels

Panoptic-DeepLab: Instances



center




center
offsets

- For instance segmentation:
 - Instance center prediction (= center of mass)
 - For each instance pixel: estimate offset vector that points towards instance center
- At inference time: Use offsets to assign instance id of closest instance center.

Panoptic DeepLab Summary

- **Shared encoder** with different decoders for semantics and instances
- **Center** and **center offsets** for bottom-up clustering of instances
- Aggregation of semantic and instance mask to produce instances

Need for labeled data



KITTI



NuScenes



Waymo Open
Dataset

- Mostly supervised training
→ need for accurately **labeled data**
- Several research datasets available
- Companies have dedicated teams to annotate data (non-public dataset)

Automotive Dataset (Detection)

Name	Year	#Categories	#Images	Data
KITTI	2012	8	15k	B,S
BDD100K	2017	10	100k	B
ApolloScape	2018	8-35	144k	B
KAIST	2018	3	9k	B
Argoverse	2019	15	22k	B
Lyft L5	2019	9	46k	B
A2D2	2019	14	12k	B
nuScenes	2019	23	40k	B,S
Waymo Open	2019	4	200k	B

Bounding Box (B), Segmentation Masks (S)

Segmentation Datasets



name	#categories	#images	task
■ MS COCO (2014)	80	118k	I,P
■ Cityscapes (2016)	30(19)	5k	S,I,P
■ Mapillary Vistas (2017)	66	25k	S,I
■ ADE20K (2017)	2,693(150)	25k	S,I

S = Semantic Seg., I = Instance Seg., P = Panoptic Seg.


Labeling data is expensive



- Labeling data is tedious and expensive
- Examples
 - Cityscapes: ~ 1.5 h per image \rightarrow 7500 h/312 days for 5k images
 - Mapillary Vistas: ~ 1.5 h per image \rightarrow 4.2 years for 25k images
 - MS COCO: 22k h (category labeling) + 10k h (instance spotting) + 26k h (instance segmentation*) \rightarrow 6.6 years
- Not included: Validation of annotations!


* 22h per 1000 segments, ~ 1.2 M instances for 80 classes

Long Tail Problem



- Class distribution is a long-tailed distribution
- Few classes (e.g., car, pedestrian) are abundantly observable
- Most classes appear in the long tail, e.g., construction vehicles, emergency vehicles, etc.

Rare Events & Situations



- Rare events that should be properly handled
- At scale, rare events are not so rare anymore

Rare Events & Situations



© Anil Polat



© dave gostisha

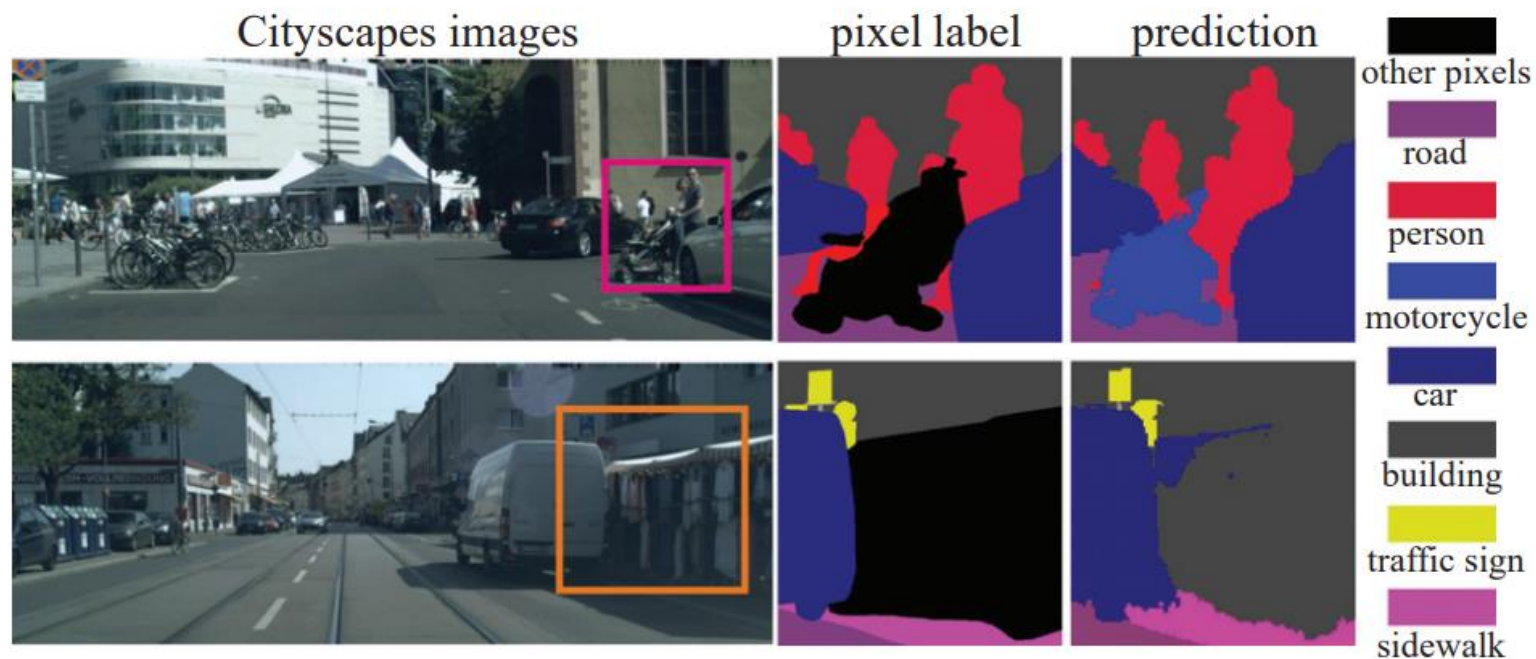


© Max Malsch




© Melanie Toomey

Open-Set Perception



- Assume that not all relevant classes are known at **training time**
- **Aim:** Learn model to identify out-of-distribution (ODD) or anomalies

OpenGAN



- Discriminator (D) distinguishes between known (closed) and unknown (open) examples
- Use generator (G) to generate additional open-set examples (besides some outlier images)

Open Challenges in Perception

- Many open challenges towards level 5 autonomous driving
 - Detection of rare events & situations
 - Challenging weather conditions
 - Adversarial attacks on perception systems
- How to deal with such situations gracefully & effectively?

Summary

- Looked at Perception stack & tasks
- Discussed CNNs as backbone of image-based perception
- Common approaches for detection, semantic segmentation and panoptic segmentation
- Challenges in visual perception and open problems

Thank you for your attention

References

- [Cordts, 2016] Cordts et al. The Cityscapes Dataset for Semantic Urban Scene Understanding, CVPR, 2016.
- [Cheng, 2020] Cheng et al. Panoptic-DeepLab: A Simple, Strong, and Fast Baseline for Bottom-up Panoptic Segmentation, CVPR, 2020.
- [Geiger,2012] Geiger et al. Are we ready for autonomous driving?, CVPR, 2012.
- [Girshick, 2014] Girshick et al. Rich Feature Hierarchies for Accurate Object Detection and Semantic Segmentation, CVPR, 2014.
- [Kirillov, 2019] Kirillov et al. Panoptic Segmentation, CVPR, 2019.
- [Kong, 2021] Kong et al. Open-Set Recognition via Open Data Generation, ICCV, 2021.
- [Law, 2018] Law et al. CornerNet: Detecting Objects as Paired Keypoints, ECCV, 2018.
- [Lin, 2014] Lin et al. Microsoft COCO: Common Objects in Context, ECCV, 2014.
- [Krizhevsky, 2012] Krizhevsky et al. ImageNet Classification with Deep Convolutional Neural Networks, NeurIPS, 2012.
- [Redmon, 2016] Redmon et al. You Only Look Once: Unified, Real-Time Object Detection, CVPR, 2016.
- [Ren, 2015] Ren et al. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks, NeurIPS, 2015.
- [Ronneberger, 2015] Ronneberger et al. U-Net: Convolutional Networks for Biomedical Image Segmentation, MICCAI, 2015.
- [Zou, 2019] Zou et al. Object Detection in 20 Years: A survey, Arxiv, 2019.