

# SELF-DRIVING CARS VIEW FROM PRACTICE



Photogrammetry & Robotics Lab

A View from Practice on Self-Driving Cars

Igor Bogoslavskyi, Apex.AI

Part of the Course: Techniques for Self-Driving Cars by  
C. Stachniss, J. Behley, N. Chebrolu, B. Mersch, I. Bogoslavskyi



BMW



Apex.AI

# AUTONOMOUS VALET PARKING

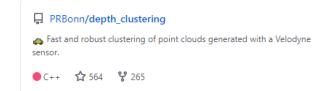


What is this lecture about?



## Open source

- More than 1000 stars on GitHub: [PRBonn/nirosus](#)



-  AUTOWARE.AUTO on Gitlab.com:  
<https://www.autoware.auto>

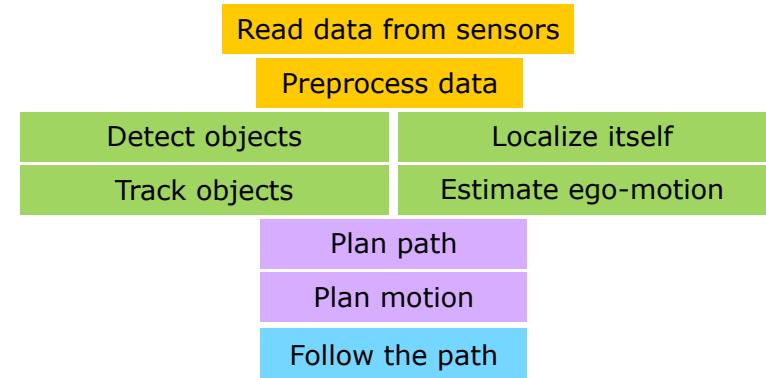


## The car must be able to



9

## The car must be able to



10



## Harder! Better! Faster! Stronger!

- Real-world problems are **harder**
- The system must be **better** integrated
- The code must run **faster**
- Safety guarantees must be **stronger**



Daft Punk by Juliana Luz

Real-world problems are **harder**



### Typical approach in academia

- Focus on publishing **novel** scientific results
- Hard to publish non-novel but robust algorithms
- **Focus is on a single part** of the autonomous driving stack, not on the full system
- Performance is measured mostly on publicly available datasets, sometimes quite old, **overfitting** to a particular one is quite common
- **Rarely tested in real world** as they require a readily-available autonomous driving stack where a component can be replaced for evaluation



15

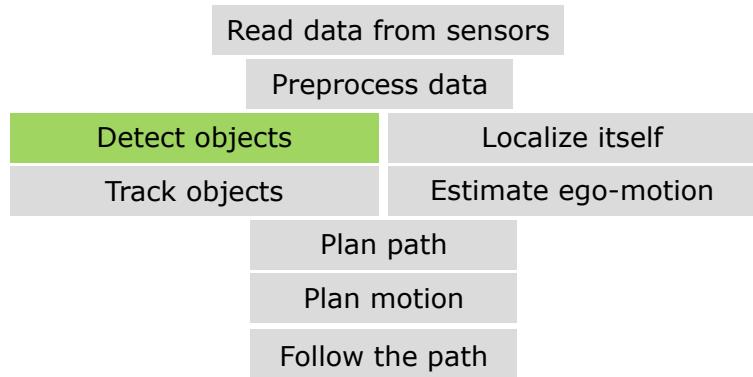
### Typical approach in industry

- Only big companies focus on **novel** scientific results
- The **simplest method** that does the job is the best
- **Focus is on the whole stack**, not on a single component
- Everything **must work in real world**
- **The tail is long**, it is not enough to nail 99% performance



16

## The car must be able to



17

## Object detection in academia

- Nowadays, **mostly learning-based** methods
- Both camera- and LiDAR-centric as well as the combination of the two
- Potentially also tracking through panoptic segmentation in combination with a Kalman Filter
- Mostly tested on some pre-recorded as well as publicly available data, e.g. KITTI dataset

18

## Is good object detection on KITTI enough?



## Dealing with strange situations

- The real world is complex
- KITTI has well-behaving road participants and no unidentifiable objects on the road
- While good performance on KITTI and alike datasets is a must, it is not enough for the real world
- Let us look at examples that are usually neglected in academic research

20

## Non-learnable objects



21



## Small objects lying on the road

- A car tire is about 20cm
- Typical LiDAR vertical resolution is 0.33°
- At 50m it is hard to detect such objects
- We can train a detector for car tires but not for all possible similar objects, e.g. squirrels, boxes etc.
- A free space classifier can be trained on images
- Sensor redundancy is still an open issue
- What can we run over?

23

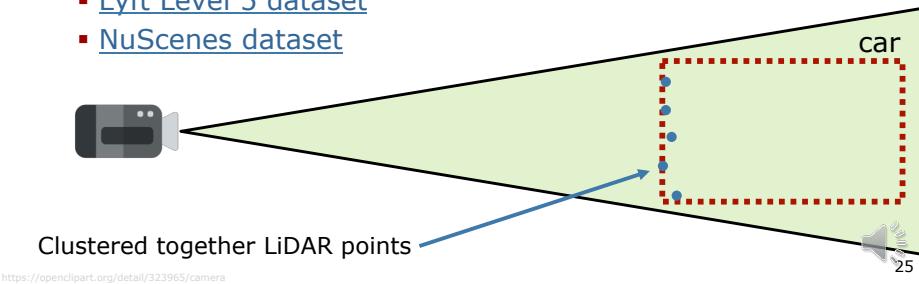
## Objects very far away (~200m)

- At 0.33° vertical LiDAR resolution there will be a 2m difference between points from different lasers
- At 0.02° horizontal LiDAR resolution there will be a 13cm distance between the points from one laser
- There will only be around 10 points on a car at this distance stemming a single line
- Driving at 50 km/h we only have 7 frames to track a car driving towards us with the same speed
- We don't want to waste even a single frame!

24

## Objects that are very far away (~200m)

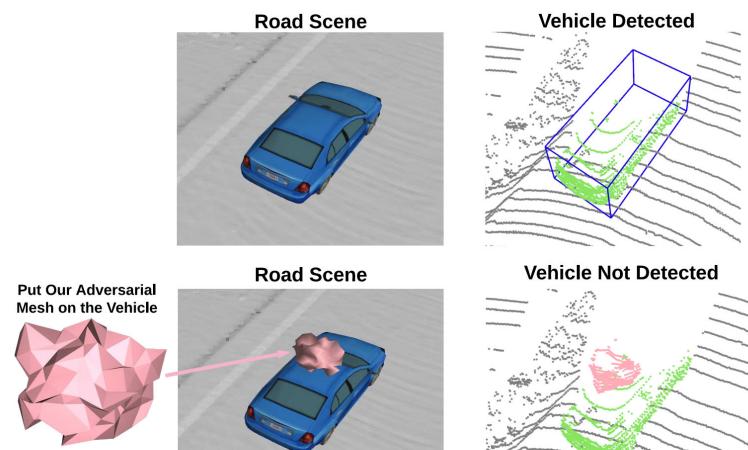
- In combination with camera detection, we can already track that object 200m away
- Datasets that offer such examples:
  - [Lyft Level 5 dataset](#)
  - [NuScenes dataset](#)



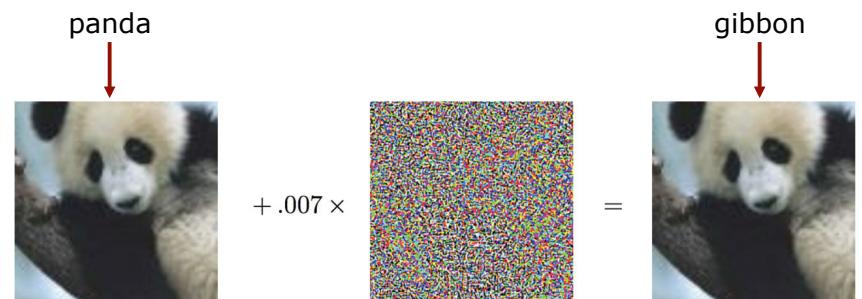
<https://openclipart.org/detail/323965/camera>

## A word on deep learning

- In recent years camera object detection has been dominated by deep learning methods
- Same happens with LiDAR object detection: deep learning methods outperform classical ones
- However, there are issues:
  - Deep learning methods mostly detect pre-learned objects
  - Prone to adversarial attacks on images and point clouds
  - Hard to predict the failures
- A fallback classical geometric approach is needed



[\[Physically Realizable Adversarial Examples for LiDAR Object Detection, J. Tu et al., CVPR 2020\]](#)



[\[Explaining and Harnessing Adversarial Examples, I. J. Goodfellow et al., ICLR 2015\]](#)



## Typical perception setup

- Camera object detection and classification
- Geometric 3D clustering, object detection and tracking, potentially including radar
- Top-down object detection on occupancy 3D grid
- Nowadays also deep learning object detection on 3D clouds directly or on range images + semantic segmentation
- Free space detection with cameras and from 3D

31

## Which sensors to use for perception?

### ▪ No single sensor is enough

- Cameras are hard to work with in the dark and don't provide metric information
- LIDAR does not work well when it is wet, snowy or when objects don't reflect light, detect phantom objects
- Radar has low resolution but measures speed and sees through fog and rain

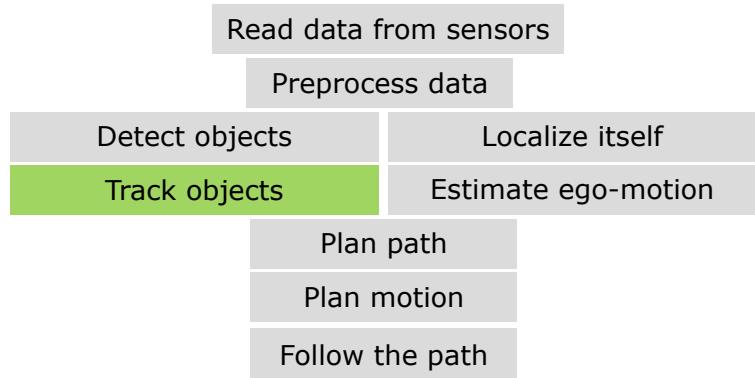
### ▪ No single method is enough

- Geometric methods don't provide enough accuracy
- Deep learning methods can fail unpredictably

30

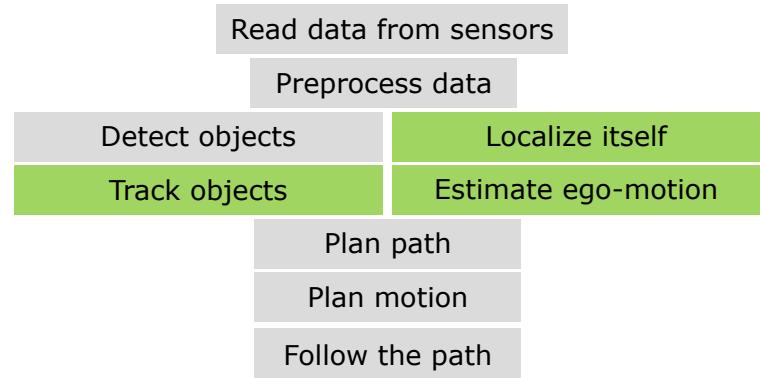


## The car must be able to



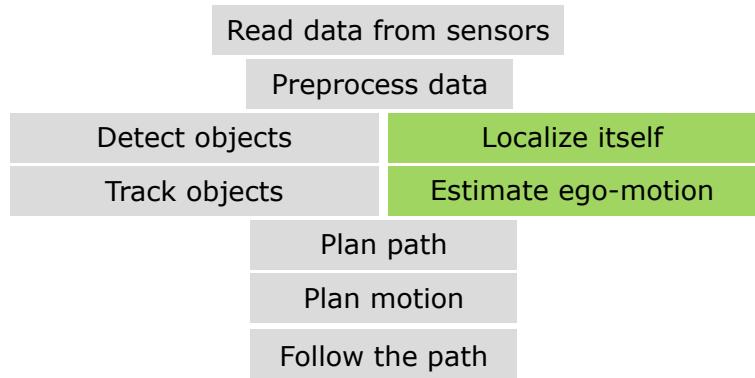
33

## The car must be able to



34

## The car must be able to



35

## Localization for self driving cars

- Most of the components of the self driving car stack rely on localization
- The car must know where it is to:
  - Estimate its own movement
  - Estimate movement of other objects
  - Plan where to drive
  - Execute the planned maneuvers
- To localize, the car needs a map

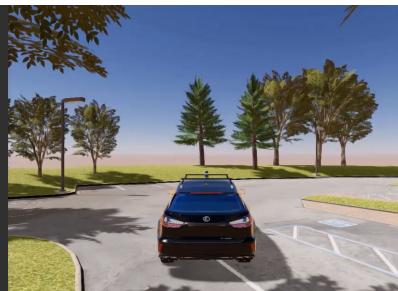
36

## Maps have different layers

1. A **3D layer** created by LiDARs
2. A **semantic layer** with information about:
  - Road lanes with driving directions and connections
  - Speed limits and other information
  - Positions of interesting objects like traffic lights
3. Optionally a **layer with dynamic changes**, e.g. road construction or roadblocks

37

## 3D Mapping



39

## Creating a 3D map layer

- Data from 3D LiDAR(s)
- Sequential registration with some flavor of Iterative Closest Point (ICP) algorithm
- Loop closure when entering already visited areas
- Pose graph optimization upon loop closures
- Stored as a collection of point clouds that can be loaded on demand

38

## Real world is much more complex

- There might be dynamic objects in the data
- Perception integration needed to detect and remove the dynamic objects
- Maintaining a full 3D map is complicated
  - Detecting changes is a non-trivial task
  - 3D maps require lots of storage
  - Big 3D maps are expensive to compute
- Online 3D map construction without the dynamics is still a non-trivial task

40

## Mapping in academia vs industry

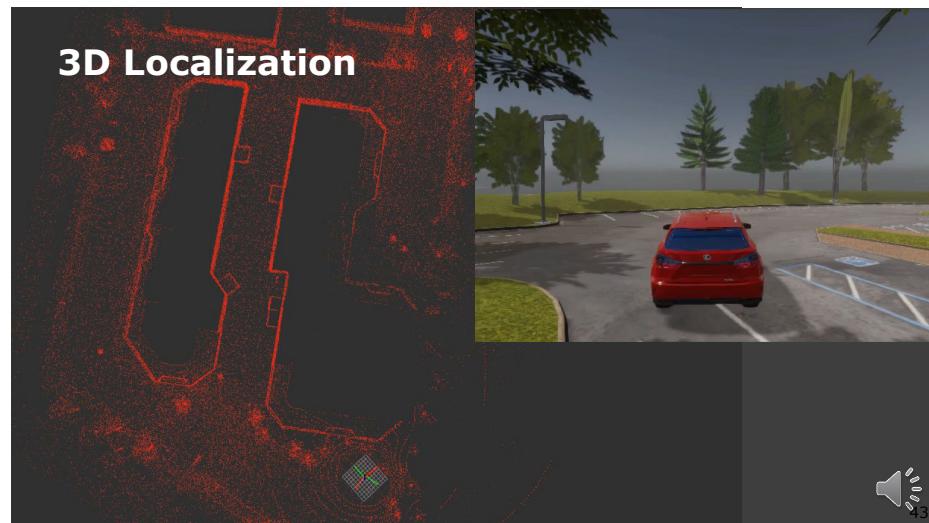
- In academia **online SLAM** is still a vibrant topic
- The papers strive to provide robust methods for map construction online
- Autonomous driving companies rarely need map construction to happen online
- Usually, **SLAM is performed offline**
- Systems are designed in a way to easily modify loop closure connections to improve the resulting map

41

## Localization in a 3D map is largely “solved”

- Data comes from 3D LiDAR(s)
- Point clouds are registered **to a given map** with some flavor of Iterative Closest Point (ICP)
- A semantic map layer can help but it is not a must
- Alternatively, localization can be done visually (not in this lecture)

42

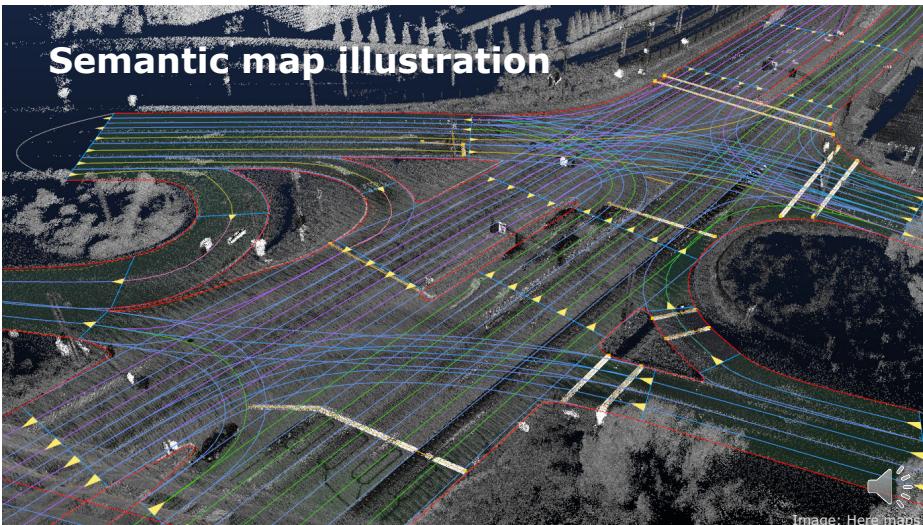


43

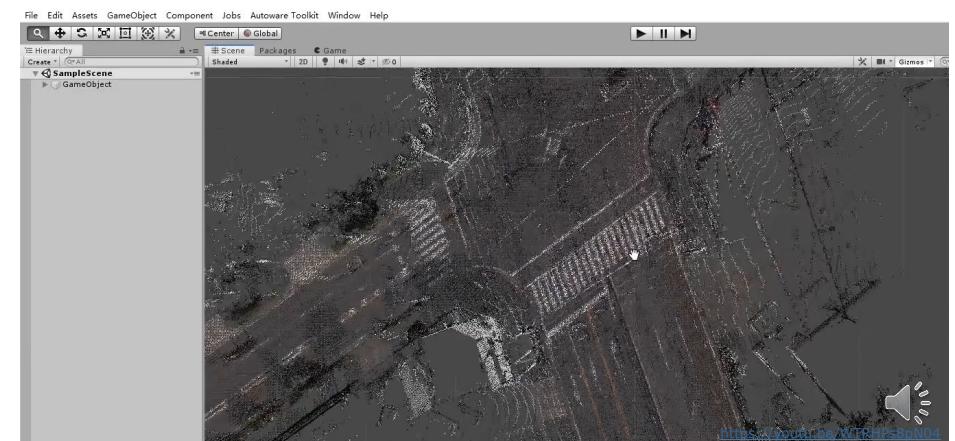
## Semantic map layer

- Contains additional information about the environment:
  - Road lanes with driving directions
  - Traffic lights
  - Speed limits
  - Additional features
- All of these are aligned to a 3D map
- Process of aligning is semi-manual

44



## Creating a semantic map is labor-intensive



## Semantic maps for traffic light detection

- Semantic map contains positions of all traffic lights
- Semantic map is aligned with the 3D map
- The car is localized in the 3D map
- This provides a good guess where the traffic lights are with respect to the car position
- This changes the traffic light detection & classification problem into pure classification

## Vanilla traffic light detection



## Semantic maps for traffic light detection



49

## What if we can't align to a 3D map?

- Maintaining a global 3D map is hard
- It is hard to scale over multiple cities/countries
- It takes a lot of processing power to align to a 3D map with full 3D point clouds
- Not all cars have a high-definition LiDAR on the roof



51

## What to do when the world changes

- Construction sites open and close, speed limits change, lanes are added and removed
- **Dealing with changes is very hard!**
- **Typical approaches:**
  - Have a constantly-deployed mapping-vehicle fleet that updates the map for all other vehicles
  - Implement change detection algorithms that allow the whole fleet to suggest changes
- Regardless of the approach manual supervision is usually needed

50

## Aligning to an HD Feature Map instead

- HD Maps are very **similar to semantic maps** but with additional features
- A decision on important information that the map must hold has to be taken before-wise
- Localization in such a map happens **based on features** extracted from the sensor data
- Sensor data is not as descriptive as with a 3D LiDAR
- Not suitable for levels 4/5 autonomy, but usually enough for levels 2/3

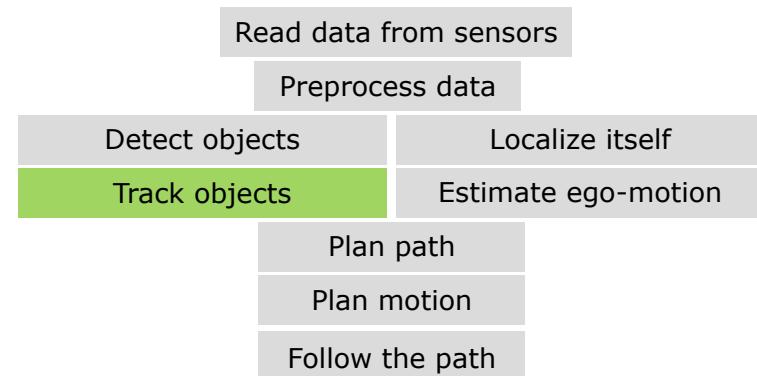
52

## Short mapping and localization summary

- Maps are usually layered
- 3D mapping and localization is mostly “solved”
- Semantic maps allow cars to know where to drive and help making some tasks easier
- The world is changing and keeping up is hard
- Not everyone can use full 3D mapping and localization, so HD feature maps are used instead
- Localization in HD feature maps is still not solved

53

## The car must be able to



54



55



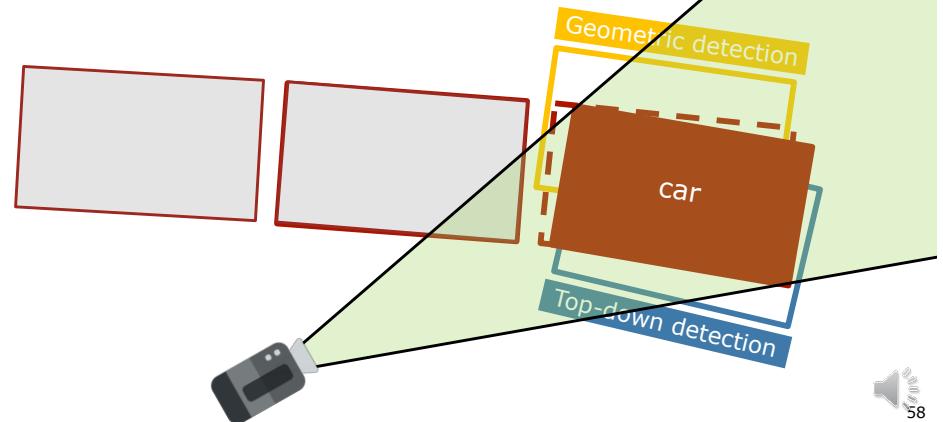
56

## Detection, tracking, and prediction pipeline

- Objects can be detected from **multiple sources**:
  - LiDAR-based geometric bounding boxes or convex hulls
  - LiDAR-based birds eye view (BEV) object detection
  - Camera object detections (provides object class)
  - Radar readings
- All of these must be assigned to each other and tracked jointly to predict what happens next
- Usually some form of **Extended Kalman Filter** or **Particle Filter** is used to track and predict
- Prediction is a huge topic

57

## Simple tracker example



58

## Tracking and prediction is hard

- A track can be hidden from the sensors for a time
- Objects can split or merge as a result of detection errors: how to assign new objects to existing ones?
- An object can change direction while unobserved

sensor

1      2

59

## Tracking and prediction is hard

- A track can be hidden from the sensors for a time
- Objects can split or merge as a result of detection errors: how to assign new objects to existing ones?
- An object can change direction while unobserved

sensor

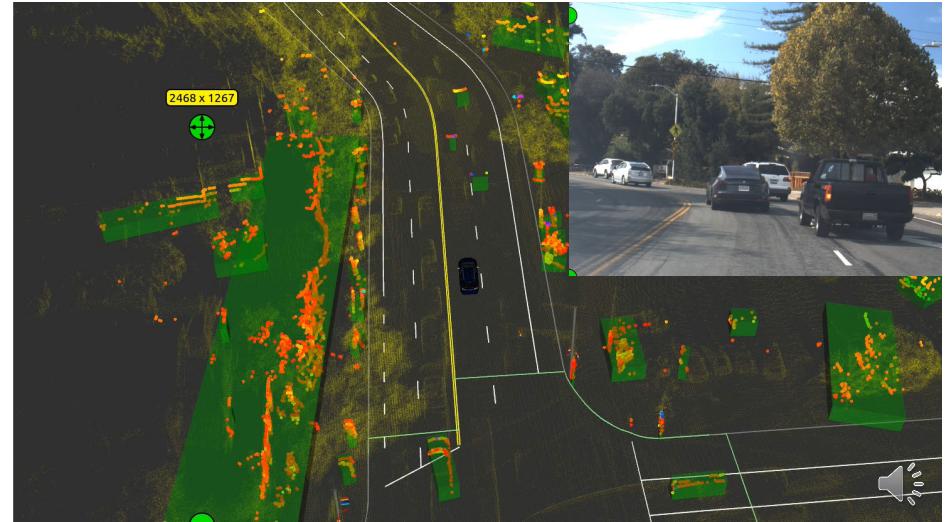
1      2

60

## Typical approach to tracking

- Standard way is to use an **Extended Kalman Filter** or a **Particle Filter** to track and predict objects
- Panoptic segmentation allows to assign an id to any object, simplifying tracking
- Prediction is still a very much open topic
- Relatively **easy to predict expected behavior** on a semantic map along lanes
- **Hard to predict unexpected behavior**

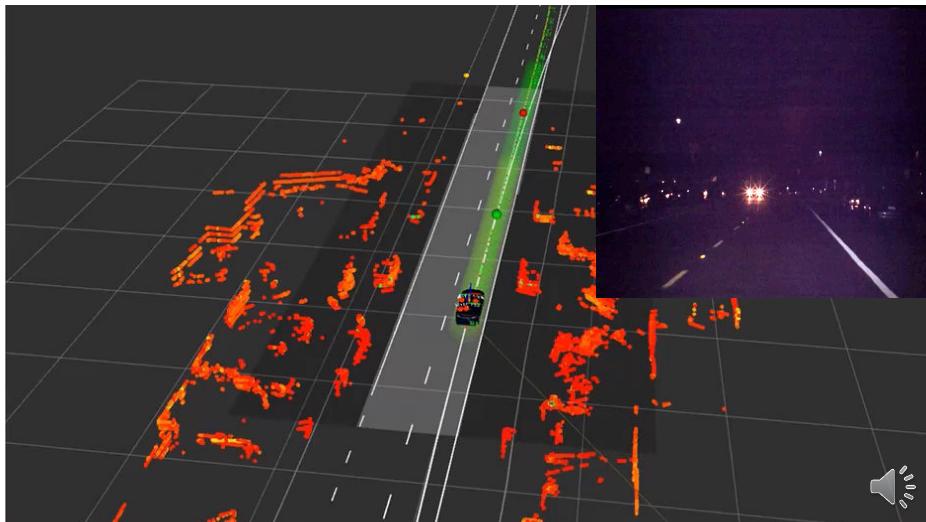
61

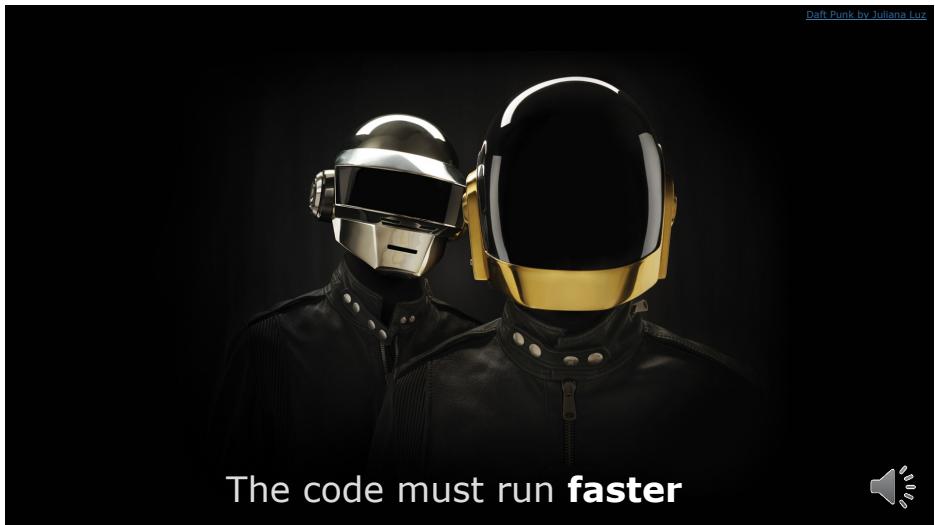


## Short summary on tracking

- Robust tracking is at the core of every self-driving car stack
- Tracking takes data from multiple sources
- Making tracking robust is still a complex task
- Predicting what the tracked objects do is even harder, especially when unexpected behavior occurs
- There is no integrated robust solution in academia

64





## Runtime – making everything run fast

- Academic research often downplays the runtime
- It is common to use “real time” in place of “at sensor frame rate” in academia
- Real time systems (RTOS) must finish every operations in a specified amount of time
- Worst-case runtime is must be **guaranteed**
- If a component takes longer than expected it might be killed which imposes **additional constraints on system design**



## Best-case scenario is not enough

- The code must fulfill performance guarantees
- Also if the processor is under load
- Also if the network is congested
- There must be performance tests present
- Careful worst-case scenario analysis is needed
- The code must be written in a way that maximizes performance by efficiently using available resources, i.e., thinking about cache locality, which work is done by which thread, priorities, etc.



**Algorithms are written and run fast!  
Are we ready to hit the streets?**





Safety guarantees must be **stronger**



Daft Punk by Juliana Lux

## Safety is extremely important!

- The car must be safe for occupants and for the surrounding people
- Safety deals with situations like:
  - What happens if a tire blew up
  - What if a sensor is unresponsive
  - What if a sensor sends bogus data
  - What if car software fails
  - How to deal with random/systematic failures
- There must be fallback options for anything that can go wrong to avoid endangering people



## How do we know the system is safe?

- It is extremely hard to build safety into the system after it is designed and implemented!
- **We design systems to be safe from the start!**
- Audit the resulting system to ensure the implementation matches the design
- The design is guided by **HARA**: Hazard Analysis and Risk Assessment



71

## ISO 26262 functional-safety standard

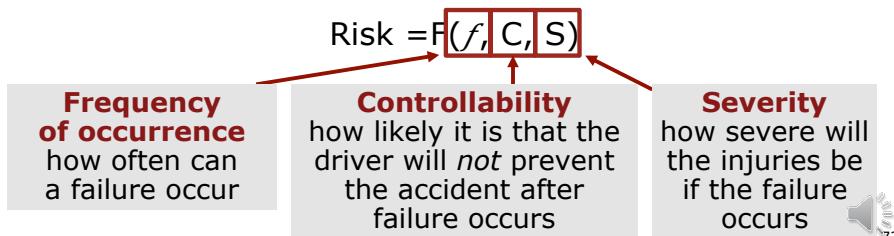
- Is an **international standard** for functional safety of road vehicles that covers:
  - **Formal definitions for errors** that might happen during the nominal operation of a car
  - **Formal definition of risk** with relation to the possible errors and ways of risk assessment and mitigation
  - **Processes to be followed** to ensure safety
- Enforces an **independent safety assessment**



72

## Risk and safety

- Risk plays the central role in safety assessment
- **Safety = absence of unreasonable risk**
- Unreasonable risk definition depends on many factors and is still widely debated



## ASIL: Automotive Safety Integrity Level

- ASIL is a function of:
  - Severity **S**: ranging from **S0** to **S3**
  - Exposure **E**: ranging from **E0** to **E4**
  - Controllability **C**: ranging from **C0** to **C3**
- ASIL levels can be **QM, A, B, C, D**
- Levels from **A** to **D** grow in safety consideration
- Table 4 of the ISO 26262-3 standard suggests an ASIL level based on the values for **S, E, C**



75

## Frequency of occurrence

- Consists of two parts:
  - **Exposure**: how often is the car in a situation where a safety hazard due to a failure can occur, denoted by **E**
  - **Failure rate**: The probability of a particular piece of hardware/software to fail
- Failure rate is **not** considered in risk assessment
- Instead, every item is assigned a so-called **ASIL** level that mitigates such risk
- ASIL level trickles down to individual components which defines how reliable they must be



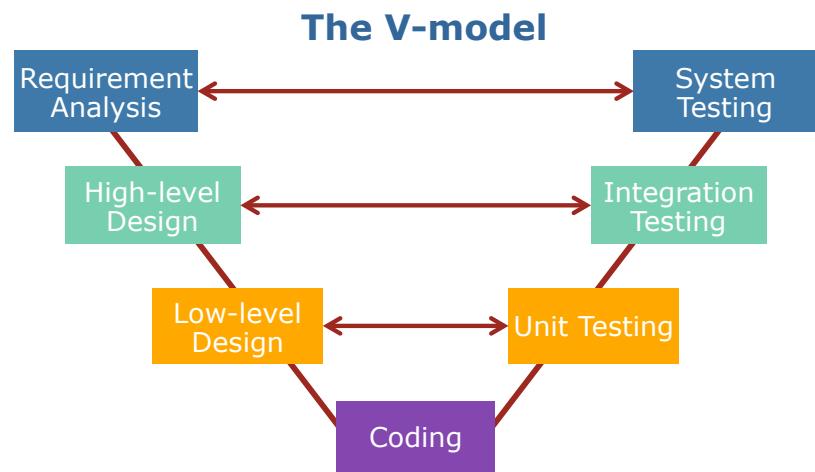
74

## Example: adaptive cruise control braking

- Adaptive highway cruise control function consists of braking, steering and propulsion systems
- Braking is an **item** in terms of ISO 26262
- This item gets an **ASIL D** level calculated from:
  - Severity **S3** – fatal injuries are likely
  - Exposure **E4** – can happen at any time on a highway
  - Controllability **C3** – complex abrupt action by the driver needed to avoid the crash
- We now derive a safety goal for the braking function of an adaptive cruise control



76



[ISO 26262-9:2018](#): Figure 1

77



- The process starts with requirement analysis on the falling edge of the diagram
- Requirement analysis corresponds to system testing on the rising edge of the diagram
- **Example requirement:**  
The car **shall** brake if a slow car merges into its lane at an intersection
- Requirement satisfaction is tested by system tests

78

## System tests

- System tests test the **behavior** of the system without knowledge of the implementation details
- These tests can be created from simulated data as well as pre-recorded real-world data
- Every test has **criteria to pass**, e.g.:
  - The car drove at least (or not more than) some distance
  - The car did not crash
  - The car reached a certain speed
  - ...

79

## Intersection system test

80



- **High level design defines the architecture**
- In our example the architectures for throttle control, object detection and prediction must be defined
- The throttle control module **shall** brake if an obstacle is predicted to be in its path
- **Integration tests** are then designed to verify the behavior of these components working together as black boxes

81



- The architecture is broken down into detailed design documents, UML diagrams etc. for each component
- A component is always a single module, class, etc.
- The **system requirements trickle down** into the single component requirements, e.g.: “the class shall have a specific API”
- Each component is tested by its **unit tests**
- **EXPECT\_EQ(42, oracle.GetMeaningOfLife());**

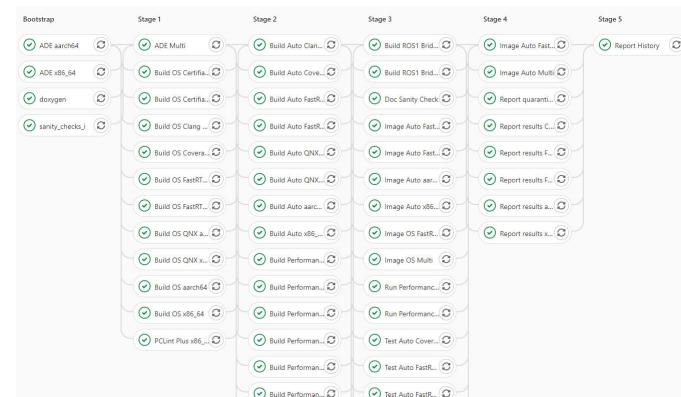
82

## Coding

- Best coding practices must be followed to avoid possible failures
- Typical coding standards: MISRA, AUTOSAR
- Static code analysis that checks these is a must
- Commits must be tagged in a specific way for traceability to a specific requirement
- A comprehensive Continuous Integration system is a must to ensure the code is always in a good state

83

## Extensive testing upon every push



84

## Languages can also be unsafe

- This course is in Python
- If you just want to work in perception by designing artificial neural networks, then it is enough
- Otherwise, C++ is currently a standard language (see [Modern C++ course by Ignacio Vizzo](#))
- Needs to have the speed, to manage resources in a deterministic and direct way and to follow guidelines for safety critical development

 85

## ISO 26262 is not the only standard

- ISO 26262 mostly targets vehicles with a driver
- Extensions and modifications are probably needed for full self driving cars
- Other standards are being developed to address specifically autonomous vehicles, e.g. [UL4600](#), [ISO/DIS 23150](#)
- Some companies design their own standards
- Expect more to happen in this field in future

 87

## Example safety consideration: memory

- Dynamically allocating memory is a blocking operation, e.g. `vector.push_back(42);`
  - It is hard to predict the runtime of this operation
  - All operations in safety-critical system must be deterministically time-bound
- **Solution:**
- Use static memory where possible
  - Allocate dynamic memory in initialization phase, only access it in the operation phase

 86

## Safety certification

- Safe software can be formally certified
- Note, that **certified != safe**
- Manufacturers of technical products must ensure that these are developed according to all aspects of functional safety and state-of-the-art science due to law of **liability**
- In case of damage, the **manufacturer must submit evidence** of the above
- Certification is a good starting point for reasoning

 88

## Safety summary

- Car must be safe to drive on public roads
- Safety is the absence of unreasonable risk derived from hazard analysis and risk assessment
- Following ISO 26262 can simplify system safety assessment and helps reducing risk of failure
- Safety certification can help in dealing with liability
- Safety enforces changes onto the system development and management process
- Safety must be a first-class citizen

89

## Closing thoughts

90

## Summary

- It is not enough to have nail 99% of the cases, see a great talk by Andrej Karpathy:  
<https://youtu.be/hx7BXih7zx8>
- Most of the work goes to handling corner cases
- Care should be taken about runtime and safety
- There is overlap but also discrepancy between industry and academia
- There is still lots of hard challenges!

91

## Where to go from here

- Start a PhD and work on relevant topics
- Contribute to open source
- Test your algorithms in challenges:
  - [Lyft Level 5 Kaggle challenge](#)
  - [NuScenes dataset and challenges](#)
  - [Waymo dataset and challenges](#)
  - Others...
- Pick people you are passionate about and try to get to work with them on relevant problems in industry

92

[Daft Punk by Juliana Luz](#)



**Our work is never over!**

